

```
In [1]: import keras
        from keras.datasets import mnist
```

Using TensorFlow backend.

```
In [2]: ((train_X, train_Y), (test_X, test_Y)) = mnist.load_data()
```

```
In [3]: print(train_X.shape)
```

(60000, 28, 28)


```
In [4]: print(test_X.shape)
```

(10000, 28, 28)

```
In [5]: import tensorflow as tf
        import matplotlib.pyplot as plt
```

```
In [6]: print(train_Y[0])
        plt.imshow(tf.squeeze(train_X[0]), cmap = 'Greys')
        plt.show()
```

5



```
In [7]: train_X = train_X.reshape(-1, 28, 28, 1)
test_X = test_X.reshape(-1, 28, 28, 1)
train_X = train_X.astype('float32')
test_X = test_X.astype('float32')
train_X = train_X / 255
test_X = test_X / 255
```

```
In [8]: from keras.utils import to_categorical
```

```
In [9]: train_Y_one_hot = to_categorical(train_Y)
test_Y_one_hot = to_categorical(test_Y)
```

```
In [10]: from keras.models import Sequential
from keras.layers import Dense, Activation, Flatten, Conv2D, MaxPooling2D
```

```
In [11]: model = Sequential()
model.add(Conv2D(64, (3, 3), input_shape = (28, 28, 1)))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size = (2, 2)))
model.add(Conv2D(64, (3, 3)))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size = (2, 2)))
model.add(Flatten())
model.add(Dense(64))
model.add(Dense(10))
model.add(Activation('softmax'))
```

```
In [12]: model.compile(loss = keras.losses.categorical_crossentropy, optimizer = keras.optimizers.Adam(), metrics = ['accuracy'])
```

```
In [13]: model.fit(train_X, train_Y_one_hot, batch_size = 64, epochs = 10)
```

```
Epoch 1/10  
60000/60000 [=====] - 63s 1ms/step - loss: 0.1468 - accuracy: 0.9567  
Epoch 2/10  
60000/60000 [=====] - 65s 1ms/step - loss: 0.0464 - accuracy: 0.9856  
Epoch 3/10  
60000/60000 [=====] - 68s 1ms/step - loss: 0.0333 - accuracy: 0.9894  
Epoch 4/10  
60000/60000 [=====] - 71s 1ms/step - loss: 0.0259 - accuracy: 0.9918  
Epoch 5/10  
60000/60000 [=====] - 67s 1ms/step - loss: 0.0203 - accuracy: 0.9939  
Epoch 6/10  
60000/60000 [=====] - 64s 1ms/step - loss: 0.0176 - accuracy: 0.9941  
Epoch 7/10  
60000/60000 [=====] - 64s 1ms/step - loss: 0.0138 - accuracy: 0.9953  
Epoch 8/10  
60000/60000 [=====] - 65s 1ms/step - loss: 0.0119 - accuracy: 0.9960  
Epoch 9/10  
60000/60000 [=====] - 66s 1ms/step - loss: 0.0098 - accuracy: 0.9968  
Epoch 10/10  
60000/60000 [=====] - 66s 1ms/step - loss: 0.0083 - accuracy: 0.9968
```

```
Out[13]: <keras.callbacks.callbacks.History at 0x1fd0215a908>
```

```
In [14]: model.save('model.h5')
```

```
In [15]: test_loss, test_acc = model.evaluate(test_X, test_Y_one_hot)
```

```
10000/10000 [=====] - 3s 274us/step
```


```
In [16]: print('Test loss:', test_loss)
```

```
Test loss: 0.04015183128897649
```

```
In [17]: print('Test accuracy:', test_acc)
```

```
Test accuracy: 0.9909999966621399
```

```
In [18]: plt.imshow(test_X[0].reshape(28, 28), cmap = plt.cm.binary)
plt.show()
```



```
In [19]: predictions = model.predict(test_X)
```

```
In [20]: import numpy as np
```

```
In [21]: print('Predicted value:', np.argmax(np.round(predictions[0])))
```

Predicted value: 7

```
In [ ]:
```